


16

Multi Threading

Objectives

After completing this lesson, you should be able to:

- Describe operating system task scheduling
- Create worker threads using `Runnable` and `Callable`
- Use an `ExecutorService` to concurrently execute tasks
- Identify potential threading problems
- Use `synchronized` and `concurrent atomic` to manage atomicity
- Use monitor locks to control the order of thread execution
- Use the `java.util.concurrent` collections



Task Scheduling

Modern operating systems use preemptive multitasking to allocate CPU time to applications. There are two types of tasks that can be scheduled for execution:

- **Processes:** A process is an area of memory that contains both code and data. A process has a thread of execution that is scheduled to receive CPU time slices.
- **Thread:** A thread is a scheduled execution of a process. Concurrent threads are possible. All threads for a process share the same data memory but may be following different paths through a code section.

Legacy Thread and Runnable

Prior to Java 5, the `Thread` class was used to create and start threads. Code to be executed by a thread is placed in a class, which does either of the following:

- Extends the `Thread` class
 - Simpler code
- Implements the `Runnable` interface
 - More flexible
 - `extends` is still free.

Extending Thread

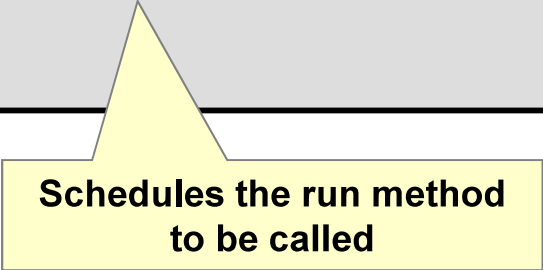
Extend `java.lang.Thread` and override the `run` method:

```
public class ExampleThread extends Thread {  
    @Override  
    public void run() {  
        for(int i = 0; i < 10; i++) {  
            System.out.println("i:" + i);  
        }  
    }  
}
```

Starting a Thread

After creating a new Thread, it must be started by calling the Thread's start method:

```
public static void main(String[] args) {  
    ExampleThread t1 = new ExampleThread();  
    t1.start();  
}
```



**Schedules the run method
to be called**

Implementing Runnable

Implement `java.lang.Runnable` and implement the `run` method:

```
public class ExampleRunnable implements Runnable {  
    private final String name;  
    public ExampleRunnable(String name) {  
        this.name = name;  
    }  
    @Override  
    public void run() {  
        for (int i = 0; i < 10; i++) {  
            System.out.println(name + ":" + i);  
        }  
    }  
}
```

Executing Runnable Instances

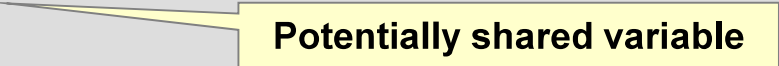
After creating a new `Runnable`, it must be passed to a `Thread` constructor. The `Thread`'s `start` method begins execution:

```
public static void main(String[] args) {  
    ExampleRunnable r1 = new ExampleRunnable();  
    Thread t1 = new Thread(r1);  
    t1.start();  
}
```


A Runnable with Shared Data

Static and instance fields are potentially shared by threads.


```
public class ExampleRunnable implements Runnable {  
    private int i;  
  
    @Override  
    public void run() {  
        for(i = 0; i < 100; i++) {  
            System.out.println("i:" + i);  
        }  
    }  
}
```



One Runnable: Multiple Threads

An object that is referenced by multiple threads can lead to instance fields being concurrently accessed.

```
public static void main(String[] args) {  
    ExampleRunnable r1 = new ExampleRunnable();  
    Thread t1 = new Thread(r1);  
    t1.start();  
    Thread t2 = new Thread(r1);  
    t2.start();  
}
```



A single Runnable instance

Detecting Interruption

Interrupting a thread is another possible way to request that a thread stop executing.


```
public class ExampleRunnable implements Runnable {  
    @Override  
    public void run() {  
        System.out.println("Thread started");  
        while(!Thread.interrupted()) {  
            // ...  
        }  
        System.out.println("Thread finishing");  
    }  
}
```

static Thread method

Interrupting a Thread

Every thread has an `interrupt()` and `isInterrupted()` method.

```
public static void main(String[] args) {  
    ExampleRunnable r1 = new ExampleRunnable();  
    Thread t1 = new Thread(r1);  
    t1.start();  
    // ...  
    t1.interrupt();  
}
```

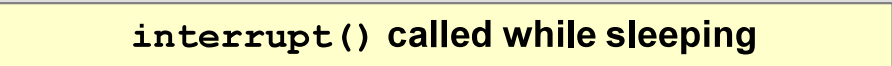


Interrupt a thread

Thread.sleep()

A Thread may pause execution for a duration of time.

```
long start = System.currentTimeMillis();
try {
    Thread.sleep(4000);
} catch (InterruptedException ex) {
    // What to do?
}
long time = System.currentTimeMillis() - start;
System.out.println("Slept for " + time + " ms");
```



interrupt() called while sleeping

Additional Thread Methods

- There are many more Thread and threading-related methods:
 - `setName(String)`, `getName()`, and `getId()`
 - `isAlive()`: Has a thread finished?
 - `isDaemon()` and `setDaemon(boolean)`: The JVM can quit while daemon threads are running.
 - `join()`: A current thread waits for another thread to finish.
 - `Thread.currentThread()`: Runnable instances can retrieve the Thread instance currently executing.
- The Object class also has methods related to threading:
 - `wait()`, `notify()`, and `notifyAll()`: Threads may go to sleep for an undetermined amount of time, waking only when the Object they waited on receives a wakeup notification.

Methods to Avoid

Some Thread methods should be avoided:

- `setPriority(int)` and `getPriority()`
 - Might not have any impact or may cause problems
- The following methods are deprecated and should never be used:
 - `destroy()`
 - `resume()`
 - `suspend()`
 - `stop()`

Deadlock

Deadlock results when two or more threads are blocked forever, waiting for each other.

```
synchronized(obj1) {  
    synchronized(obj2) {  
    }  
}
```

Thread 1 pauses after locking obj1's monitor.

```
synchronized(obj2) {  
    synchronized(obj1) {  
    }  
}
```

Thread 2 pauses after locking obj2's monitor.

The `java.util.concurrent` Package

Java 5 introduced the `java.util.concurrent` package, which contains classes that are useful in concurrent programming. Features include:

- Concurrent collections
- Synchronization and locking alternatives
- Thread pools
 - Fixed and dynamic thread count pools available
 - Parallel divide and conquer (Fork-Join) new in Java 7

Recommended Threading Classes

Traditional `Thread` related APIs are difficult to code properly. Recommended concurrency classes include:

- `java.util.concurrent.ExecutorService`, a higher level mechanism used to execute tasks
 - It may create and reuse `Thread` objects for you.
 - It allows you to submit work and check on the results in the future.
- The Fork-Join framework, a specialized work-stealing `ExecutorService` new in Java 7

java.util.concurrent.ExecutorSe

An `ExecutorService` is used to execute tasks.

- It eliminates the need to manually create and manage threads.
- Tasks **might** be executed in parallel depending on the `ExecutorService` implementation.
- Tasks can be:
 - `java.lang.Runnable`
 - `java.util.concurrent.Callable`
- Implementing instances can be obtained with `Executors`.

```
ExecutorService es = Executors.newCachedThreadPool();
```

Example ExecutorService

This example illustrates using an `ExecutorService` to execute `Runnable` tasks:

```
package com.example;
import java.util.concurrent.ExecutorService;
import java.util.concurrent.Executors;
public class ExecutorExample {
    public static void main(String[] args) {
        ExecutorService es = Executors.newCachedThreadPool();
        es.execute(new ExampleRunnable("one"));
        es.execute(new ExampleRunnable("two"));
        es.shutdown();
    }
}
```

**Execute this Runnable task
sometime in the future**

Shut down the executor

Shutting Down an ExecutorService

Shutting down an `ExecutorService` is important because its threads are nondaemon threads and will keep your JVM from shutting down.

```
es.shutdown();  
try {  
    es.awaitTermination(5, TimeUnit.SECONDS);  
} catch (InterruptedException ex) {  
    System.out.println("Stopped waiting early");  
}
```

**Stop accepting new
Callables.**

**If you want to wait for the Callables to
finish**

java.util.concurrent.Callable

The Callable interface:

- Defines a task submitted to an `ExecutorService`
- Is similar in nature to `Runnable`, but can:
 - Return a result using generics
 - Throw a checked exception

```
package java.util.concurrent;  
  
public interface Callable<V> {  
    V call() throws Exception;  
}
```

Example Callable Task

```
public class ExampleCallable implements Callable {  
  
    private final String name;  
    private final int len;  
    private int sum = 0;  
  
    public ExampleCallable(String name, int len) {  
        this.name = name;  
        this.len = len;  
    }  
  
    @Override  
    public String call() throws Exception {  
        for (int i = 0; i < len; i++) {  
            System.out.println(name + ":" + i);  
            sum += i;  
        }  
        return "sum: " + sum;  
    }  
}
```

Return a String from this task: the sum of the series

java.util.concurrent.Future

The Future interface is used to obtain the results from a Callable's V call() method.

```
Future<V> future = es.submit(callable);  
//submit many callables  
try {  
    V result = future.get();  
} catch (ExecutionException|InterruptedException ex) {  
}
```

ExecutorService controls when the work is done.

Gets the result of the Callable's call method (blocks if needed).

If the Callable threw an Exception

Example

```
public static void main(String[] args) {  
  
    ExecutorService es = Executors.newFixedThreadPool(4);  
    Future<String> f1 = es.submit(new ExampleCallable("one",10));  
    Future<String> f2 = es.submit(new ExampleCallable("two",20));  
  
    try {  
        es.shutdown();  
        es.awaitTermination(5, TimeUnit.SECONDS);  
        String result1 = f1.get();  
        System.out.println("Result of one: " + result1);  
        String result2 = f2.get();  
        System.out.println("Result of two: " + result2);  
    } catch (ExecutionException | InterruptedException ex) {  
        System.out.println("Exception: " + ex);  
    }  
}
```

Wait 5 seconds for the tasks to complete

Get the results of tasks f1 and f2

Threading Concerns

- Thread Safety
 - Classes should continue to behave correctly when accessed from multiple threads.
- Performance: Deadlock and livelock
 - Threads typically interact with other threads. As more threads are introduced into an application, the possibility exists that threads will reach a point where they cannot continue.

Shared Data

Static and instance fields are potentially shared by threads.

```
public class SharedValue {  
    private int i;  
  
    // Return a unique value  
    public int getNext() {  
        return i++;  
    }  
}
```

Potentially shared variable

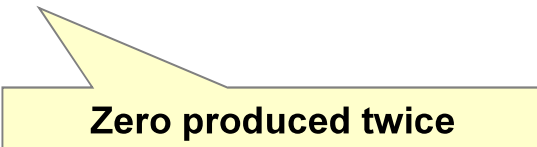
Problems with Shared Data

Shared data must be accessed cautiously. Instance and static fields:

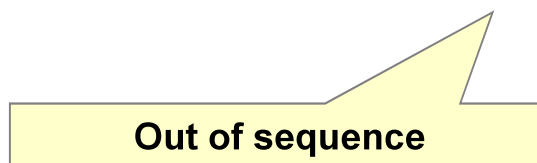
- Are created in an area of memory known as heap space
- Can potentially be shared by any thread
- Might be changed concurrently by multiple threads
 - There are no compiler or IDE warnings.
 - “Safely” accessing shared fields is your responsibility.

Two threads accessing an instance of the `SharedValue` class might produce the following:

`i:0,i:0,i:1,i:2,i:3,i:4,i:5,i:6,i:7,i:8,i:9,i:10,i:12,i:11 ...`



Zero produced twice



Out of sequence

Nonshared Data

Some variable types are never shared. The following types are always thread-safe:

- Local variables
- Method parameters
- Exception handler parameters
- Immutable data

The `synchronized` Keyword

The `synchronized` keyword is used to create thread-safe code blocks. A `synchronized` code block:

- Causes a thread to write all of its changes to main memory when the end of the block is reached
- Is used to group blocks of code for exclusive execution
 - Threads block until they can get exclusive access
 - Solves the atomic problem

synchronized Methods

```
3 public class SynchronizedCounter {  
4     private static int i = 0;  
5  
6     public synchronized void increment() {  
7         i++;  
8     }  
9  
10    public synchronized void decrement() {  
11        i--;  
12    }  
13    14    public synchronized int getValue() {  
15        return i;  
16    }  
17 }
```

synchronized Blocks

```
18  public void run() {  
19      for (int i = 0; i < countSize; i++) {  
20          synchronized(this) {  
21              count.increment();  
22              System.out.println(threadName  
23                  + " Current Count: " + count.getValue());  
24          }  
25      }  
26  }
```


Object Monitor Locking

Each object in Java is associated with a monitor, which a thread can lock or unlock.

- `synchronized` methods use the monitor for the `this` object.
- `static synchronized` methods use the classes' monitor.
- `synchronized` blocks must specify which object's monitor to lock or unlock.

```
synchronized ( this ) { }
```

- `synchronized` blocks can be nested.

Threading Performance

To execute a program as quickly as possible, you must avoid performance bottlenecks. Some of these bottlenecks are:

- Resource Contention: Two or more tasks waiting for exclusive use of a resource
- Blocking I/O operations: Doing nothing while waiting for disk or network data transfers
- Underutilization of CPUs: A single-threaded application uses only a single CPU

Performance Issue: Examples

- **Deadlock** results when two or more threads are blocked forever, waiting for each other.

```
synchronized(obj1) {  
    synchronized(obj2) {  
    }  
}
```

Thread 1 pauses after locking obj1's monitor.

```
synchronized(obj2) {  
    synchronized(obj1) {  
    }  
}
```

Thread 2 pauses after locking obj2's monitor.

- **Starvation and Livelock**

Summary

In this lesson, you should have learned how to:

- Describe operating system task scheduling
- Use an `ExecutorService` to concurrently execute tasks
- Identify potential threading problems
- Use `synchronized` and `concurrent atomic` to manage atomicity
- Use monitor locks to control the order of thread execution
- Use the `java.util.concurrent` collections

